

A semantic news aggregator in Python using Dbpedia, Cubicweb and Scikits-learn

Vincent Michel - Logilab

26 août 2011

- 1 Introduction
- 2 Extracting information from RSS news
- 3 Analyzing RSS news


With a bunch of RSS news...

- *Google Borrows Apple Strategy : Google's deal underscores the allure of a business model pioneered...*
- *Japan disaster plant cold shutdown could face delay : TOKYO (Reuters) - Tokyo Electric Power Co said on Wednesday...*
- *Libya shows signs of slipping from Muammar Gaddafi's grasp : Supply lines to capital in peril as coastal cities fall...*

... how can we analyze them in Python ?

- Clustering (grouping) RSS (e.g. *Google News*).
- **Extracting/synthetizing information.**
- **Providing useful/original visualisation and analytics tools.**

Semantic ?



“Linking Open Data cloud diagram, by Richard Cyganiak and Anja Jentzsch.

<http://lod-cloud.net/>

Fetching, storing and querying the data → CubicWeb (*)

- ✓ Semantic CMS, high-level database management with metadata.
- ✓ Multiple sources : RSS, micro-blogging, SQL database, ...
- ✓ Based on **PostgreSQL** : deals with (very) large database.

Data-mining and machine learning → Scikits-learn (*)

- ✓ Easy-to-use and general-purpose machine learning in Python.
- ✓ Unsupervised learning, supervised learning, model selection ...

Semantic information database → Dbpedia (*)

- ✓ $\sim 8.10^6$ articles with *abstracts*, *images*, ... from Wikipedia.
- ✓ $\sim 0,6.10^6$ categories, 273 types (*e.g. person, place, ...*)
- ✓ $\sim 100.10^6$ links between articles, categories and types.

(*) Open Source/Creative Commons !!!

Storing RSS information with **Cubicweb**

Each object (or **entity**) is defined in a **schema** and may be displayed using different **views**.

Storing RSS

Define (in *schema.py*) how RSS should be stored in the database.

```
class RSSArticle(EntityType):
    title = String() # Title of the feed
    uri = String(unique=True) # Uri of the rss feed
    content = String() # Content of the feed
```

Fetching RSS (based on *feedparser* and *BeautifulSoup*)

Simply construct a **source**, by giving an **URL** and a **parser** :

```
url = u'http://feeds.bbc1.co.uk/news/world/rss.xml'
s = session.create_entity('CWSource', name=u'BBCNews-World', url=url,
                           type=u'datafeed', parser=u'rss-parser',
                           config=u'synchronization-interval=240min')
s.pull_data(session)
```

7 english/american journals (*The New York Times*, ...)

Storing Dbpedia page

Define (in *schema.py*) how Dbpedia pages should be stored in the database.

```
class DbpediaPage(EntityType):
    uri = String(unique=True, indexed=True) # Uri of the ressource
    label = String(indexed=True) # http://www.w3.org/2000/01/rdf-schema
    pageid = String() # http://dbpedia.org/ontology/wikiPageID
    abstract = String() # http://dbpedia.org/ontology/abstract
    homepage = String() # http://xmlns.com/foaf/0.1/homepage
    thumbnail = String() # http://dbpedia.org/ontology/thumbnail
    depiction = String() # http://xmlns.com/foaf/0.1/depiction
    wikipedia = String() # http://xmlns.com/foaf/0.1/page
    latitude = String() # http://www.w3.org/2003/01/geo/wgs84_pos
    longitude = String() # http://www.w3.org/2003/01/geo/wgs84_pos
```

Storing all dbpedia information ($\sim 9.10^6$ pages, $\sim 100.10^6$ links, 20Go) in Cubicweb takes **less than 24 hours**.

→ See the full schema

What can we do with the RSS news stored in database ?

- 1 extract **relevant features of the data**.
- 2 construct a **usable (i.e. matrix) representation of the data**.
- 3 **clusters** (group) RSS together.
- 4 deeper **analyze/visualization** of the information.

Example sentence :

“Google is to buy mobile phone manufacturer Motorola Mobility, allowing it to mount a serious challenge to Apple Inc.”

- 1 Introduction
- 2 Extracting information from RSS news**
- 3 Analyzing RSS news

Extracting information : Classical approaches (**Scikits-learn**)

Char-N-gram

Extracts tokens of N characters from a text.

```
from scikits.learn.feature_extraction.text import CharNGramAnalyzer
analyzer = CharNGramAnalyzer(min_n=3, max_n=6)
tokens = analyzer.analyze(sentence)
```

450 tokens : *'goo', 'oog', 'ogl', ..., 'mob', 'obi', 'bil', ...*

Word-N-gram

Extracts tokens of N words from a text.

```
from scikits.learn.feature_extraction.text import WordNGramAnalyzer
analyzer = WordNGramAnalyzer(min_n=3, max_n=6)
tokens = analyzer.analyze(sentence)
```

58 tokens : *'google is to', 'is to buy', ..., 'serious challenge to', ...*

- ✗ Many irrelevant features.
- ✗ Features do not carry lots of contextual information (i.e. **understandable by humans**).

Main Hypothesis : Only things that exist in Dbpedia (i.e Wikipedia) have some interest in news analysis → Named Entities Recognition (NER)

Dbpedia feature extraction (**Cubicweb/Dbpedia**)

```
from cubes.semnews.views.nertools import DbpediaEntitiesAnalyzer
analyzer = DbpediaEntitiesAnalyzer(session, lang='en')
tokens = analyzer.analyze(sentence)
```

→ 3 tokens : 'Apple Inc', 'Google', 'Motorola'

<ENAMEX TYPE="ORGANIZATION"> Google</ENAMEX> is to buy mobile phone manufacturer <ENAMEX TYPE="ORGANIZATION"> Motorola</ENAMEX> Mobility, allowing it to mount a serious challenge to <ENAMEX TYPE="ORGANIZATION"> Apple Inc</ENAMEX>

→ Try it !

"DBpedia Spotlight : Shedding Light on the Web of Documents", Pablo N. Mendes et al., I-Semantics 2011

"Learning Named Entity Recognition from Wikipedia", Joel Nothman 2008

"Large-Scale Named Entity Disambiguation Based on Wikipedia Data", Silviu Cucerzan 2007

Efficient and robust feature extraction

- **Keep the meaning of a text** → interpretable features.

✗ *'... said former soldier Larry, ...'*

✗ *'... said student Larry, ...'*

✓ *'... said Larry Page, ...'*

- **Robust features based on redirections.**

e.g. 'Obama', 'Barak Obamba', 'Pres. Obama' redirects to 'Barack Obama'

Simple RQL (*Relation Query Language*) queries

- **Fast**, based on indexed SQL tables and regular expressions :

```
rset = rql('Any E WHERE E is DbpediaPage,  
          E label %(token)s', {'token': token})
```

e.g. 19 entities extracted in 4s in 765 words, among $\sim 8.10^6$ dbpedia entries.

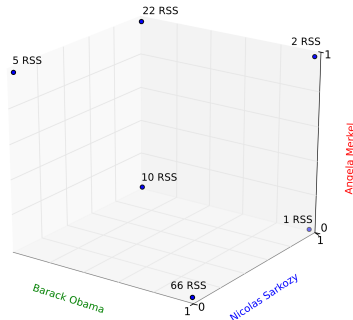

- Different labels but same URI → **cross-language feature extraction** : *e.g. Grenada/Grenade → <http://dbpedia.org/resource/Grenada>*

- 1 Introduction
- 2 Extracting information from RSS news
- 3 Analyzing RSS news**

Feature extraction : Matrix representation and storage

$$\begin{pmatrix} \text{"Obama's approval ..."} \\ \text{"Protest as Spain ..."} \\ \dots \\ \text{"Libya rebels fight ..."} \\ \text{"Obama in Spain ..."} \end{pmatrix} \rightarrow \begin{pmatrix} 0 & 0 & \dots & 1 & 0 \\ 0 & 0 & \dots & 0 & 1 \\ \dots & & & & \\ 0 & 1 & \dots & 0 & 0 \\ 0 & 0 & \dots & 1 & 1 \end{pmatrix}$$

E.g. The Obama-Merkel-Sarkozy space ...



Results stored in a **relation** (*appears in rss*) in **Cubicweb** :

```
rql('Any X WHERE X appears_in_rss Y')
```

Creating clusters (groups) of news, using the matrix representation of the data.

Meanshift algorithm (**Scikits-learn**)

- Based on locating the maxima of a density function.
- **Automatically tunes the number of clusters.**

```
from scikits.learn.cluster import MeanShift, estimate_bandwidth
bandwidth = estimate_bandwidth(X, quantile=0.005)
clustering = MeanShift(bandwidth=bandwidth)
clustering.fit(X)
labels = clustering.labels_
```

"Mean shift, mode seeking, and clustering.", Yizong Cheng. IEEE Transactions on Pattern Analysis and Machine Intelligence 1995

→ Try it !

Other possible alternatives : Ward's algorithm, K-means, ...

Exploiting information : Queries and Views

Each query returns a **result set (rset)**. A view is called on a **result set**
→ define the representation rules.

Defining a view (short version)

```
class MyEntitiesView(View):
    __regid__ = 'example-view'

    def call(self):
        rset = self.cw_rset
        for entity in rset.entities():
            do_whatever_you_want...
```

... you just have to plug a rset from a query :

```
rset = rql('Any X WHERE ...')
self.wview('example-view', rset)
```

or within an url :

`http://myapplication/?rql=Any X WHERE&vid=example-view`

A new approach for querying information from RSS!

All musical artists in the news

```
rql('DISTINCT Any E, R WHERE E appears_in_rss R,  
    E has_type T, T label "musical artist"')
```

All living office holder persons in the news

```
rql('DISTINCT Any E WHERE E appears_in_rss R,  
    E has_type T, T label "office holder",  
    E has_subject C, C label "Living people"')
```

All news that talk about Barack Obama and any scientist

```
rql('DISTINCT Any R WHERE E1 label "Barack Obama",  
    E1 appears_in_rss R, E2 appears_in_rss R,  
    E2 has_type T, T label "scientist"')
```

All news that talk about a drug

```
rql('Any X, R WHERE X appears_in_rss R,  
    X has_type T, T label "drug"')
```

Try it! ...with an xml view ...or a thumbnail view

View

```
class EntityMapView(View):  
    __regid__ = 'map'  
  
    def call(self):  
        rset = self.cw_rset  
        self.init_map()  
        for entity in rset.entities():  
            self.add_marker(entity.latitude, entity.longitude,  
                             entity.dc_title)  
        self.center_and_zoom(0, 0, 1.5)  
        self.finish_map()
```

Based on mapstraction (javascript) : <http://mapstraction.com/>

✓ **Automatically locate information from RSS news.**

Try it !

Cubicweb and Scikits-learn :

- Efficient and easy-to-use tools for data storing, querying and mining.
- Easy to plug together, only **Python** tools, all **Open Source**.

Using Dbpedia allows to extract very few highly relevant features

- ✓ Decrease the dimensionality of the data.
- ✓ Link features to millions of pages of information.

A new semantic way for querying information

- ✓ Simple information queries using RQL expressions.
- ✓ Use Dbpedia types and categories to refine the selection.

Named Entities Recognition

- use disambiguations links / more refined Regular expressions.
- add new databases (*MusicBrainz*, *diseasome*, ...).
- closely follow Wikipedia with **Dbpedia live update** :

Motorola Mobility (11 :46, 15 August 2011) 'On the 15 August, Google announced that it agreed to acquire the company.'

RSS news analyzing

- explore new algorithms : bi-clustering, ...
- add new data sources : Twitters, Blogs, ...
- *'from scikits.learn predict __future__ ...'* → use **Matrix Completion** to predict new edges in the correlation graph.

Thanks you for your attention !

Questions ?